

Assignment 7

CS 532: Introduction to Web Science

Spring 2017

Grant Atkins

Finished on April 6, 2017

1

Question

1. Find 3 users who are closest to you in terms of age, gender, and occupation. For each of those 3 users:

- what are their top 3 favorite films?
- bottom 3 least favorite films?

Based on the movie values in those 6 tables (3 users X (favorite + least)), choose a user that you feel is most like you. Feel free to note any outliers (e.g., "I mostly identify with user 123, except I did not like 'Ghost' at all").

This user is the "substitute you".

Answer

First approaching this problem I decided to use python 3.6 because the code written in the Programming Collective Intelligence book, by Toby Segaran, was also python [2]. Instead of manually picking out users from the data file provided, I wrote a script called **substituteMe.py**, shown in Listing 1, which filters all users by the gender Male “M”, the occupation of “programmer” and the age range of greater than 20 and less than 23. This script found multiple users ages 21 and 23, but I decided to go with users with age 21 since I was just recently 21. This left me with 3 users with the ids: 603, 671 and 868.

To find their favorite and least favorite movies I added a function called *findMoviesMerge*, which matched each user’s review to their movie names and return the bottom and top movies sorted by their ratings, there were of course other movies with rating 5 but I simply took the top and bottom 3 provided. Their tables for top 3 films and bottom 3 films are shown in Table 1, Table 2 and Table 3 respectively.

I identified that user 868 was the “substitute me”. This user had excellent movie choice with me liking all of his top favorite movies as well as me also disliking the same movies. Who would make a live action version of Super Mario Bros?

Rank	Top Favorite Movies	Rating	Least Favorite Movie	Rating
1	Star Wars (1977)	5.0	Platoon (1986)	1.0
2	Blade Runner (1982)	5.0	Heat (1995)	1.0
3	Twelve Monkeys (1995)	5.0	Platoon (1986)	2.0

Table 1: User 603’s favorite and least favorite movies

Rank	Top Favorite Movies	Rating	Least Favorite Movie	Rating
1	My Best Friend’s Wedding (1997)	5.0	Cop Land (1997)	1.0
2	Walk in the Clouds, A (1995)	5.0	Long Kiss Goodnight, The (1996)	1.0
3	Terminator, The (1984)	5.0	Star Trek: First Contact (1996)	1.0

Table 2: User 671’s favorite and least favorite movies

Rank	Top Favorite Movies	Rating	Least Favorite Movie	Rating
1	2001: A Space Odyssey (1968)	5.0	Lassie (1994)	1.0
2	Raiders of the Lost Ark (1981)	5.0	Super Mario Bros. (1993)	1.0
3	Empire Strikes Back, The (1980)	5.0	Herbie Rides Again (1974)	1.0

Table 3: User 868’s favorite and least favorite movies

```

1 import csv
2 from pprint import pprint as pp
3
4
5 def chooseUsers():
6     usersChosen = []
7     with open("data/u.user") as f:
8         reader = csv.reader(f, delimiter='|')
9
10        for i in reader:
11            age = int(i[1])
12            # filter parameters
13            if(i[2] == 'M' and i[3] == 'programmer' and
14               (age > 20 and age < 23)):
15                usersChosen.append(i)
16
17    return usersChosen
18
19
20 def findReviews(userIds):
21     # pairs are user id -> array of reviews
22     reviewDict = {}
23     for i in userIds:
24         reviewDict[i] = []
25     with open("data/u.data", 'r') as f:
26
27         for line in f:
28             spl = line.split()
29             for i in userIds:
30                 if(spl[0] == i):
31                     reviewDict[i].append(spl)
32
33     return reviewDict
34
35
36 def findMovie(movieId):
37     with open("data/u.item", 'r') as f:
38         reader = csv.reader(f, delimiter='|')
39         for i in reader:
40             itemId = i[0]
41             if movieId == itemId:

```

```

42         # id, name, URI
43         return (i[0], i[1], i[4])
44
45
46 def findMoviesMerge(reviewDict):
47     userMovieDict = {}
48     for userId, reviews in reviewDict.items():
49
50         userMovieDict[userId] = {}
51         moviesReviewed = []
52         botMovies = []
53         topMovies = []
54         for r in reviews:
55             movieId = r[1]
56             rating = r[2]
57
58             movie = findMovie(movieId)
59             movie = tuple(rating) + movie
60             moviesReviewed.append(movie)
61
62             # botMovies.sort(key=lambda tup: tup[0])
63             moviesReviewed.sort(key=lambda tup: tup[0])
64             botMovies = moviesReviewed[:3]
65             topMovies = moviesReviewed[-3:]
66             userMovieDict[userId]["bottomMovies"] = botMovies
67             userMovieDict[userId]["topMovies"] = topMovies
68
69     return userMovieDict
70
71
72 if __name__ == "__main__":
73     chosenUsers = chooseUsers()
74     userIds = []
75     for i in chosenUsers:
76         userIds.append(i[0])
77
78     reviewDict = findReviews(userIds)
79     with open("data/closestUsers.txt", 'w') as f:
80         pp(findMoviesMerge(reviewDict), stream=f)

```

Listing 1: Python script for determining closest 3 users

2

Question

2. Which 5 users are most correlated to the substitute you? Which 5 users are least correlated (i.e., negative correlation)?

Answer

This question's answer relied heavily off the source code provided by the Programming Collective Intelligence book [2]. I created a script called **correlateUsers.py**, shown in Listing 2, which is also used later in questions 3 and 4. This question used the *loadMovieLens*, which generated the user preferences, and *findCorrelations* methods which finds the Sim Pearson correlation coefficient between “substitute me,” user 868, and every other user based on their preferences. The results of this are shown in Tables 4 and 5 and were saved to **correlatedUsers.txt** in my Github repository [1].

User ID	Correlation
853	+1.0
857	+1.0
898	+1.0
625	+1.0
724	+1.0

Table 4: Most correlated users

User ID	Correlation
36	-1.0
404	-1.0
599	-1.0
628	-1.0
736	-0.9045340337332909

Table 5: Least correlated users

```

1
2 import csv
3 from math import sqrt
4 from pprint import pprint as pp
5
6
7 def sim_pearson(prefs, p1, p2):
8     '''
9     Returns the Pearson correlation coefficient for p1 and p2.
10    '''
11
12    # Get the list of mutually rated items
13    si = {}
14    for item in prefs[p1]:
15        if item in prefs[p2]:
16            si[item] = 1
17    # If they are no ratings in common, return 0
18    if len(si) == 0:
19        return 0
20    # Sum calculations
21    n = len(si)
22    # Sums of all the preferences
23    sum1 = sum([prefs[p1][it] for it in si])
24    sum2 = sum([prefs[p2][it] for it in si])
25    # Sums of the squares
26    sum1Sq = sum([pow(prefs[p1][it], 2) for it in si])
27    sum2Sq = sum([pow(prefs[p2][it], 2) for it in si])
28    # Sum of the products
29    pSum = sum([prefs[p1][it] * prefs[p2][it] for it in si])
30    # Calculate r (Pearson score)
31    num = pSum - sum1 * sum2 / n
32    den = sqrt((sum1Sq - pow(sum1, 2) / n) * (sum2Sq - pow(sum2,
33        2) / n))
34    if den == 0:
35        return 0
36    r = num / den
37    return r
38
39 def findCorrelations(prefs):
40     most_correlated = []
41     least_correlated = []
42     correlations = {}
43     substituteMe = str(868)
44
45     users = {}
46     for line in open('data/u.user'):
47         (user, age, gender, job, zipcode) = line.split('|')

```

```

48         users.setdefault(user, {})
49         users[user] = {'age': age, 'gender': gender,
50                        'job': job, 'zipcode': zipcode}
51
52     for user, rest in users.items():
53         if substituteMe == user:
54             pass
55         else:
56             r = sim_pearson(prefs, substituteMe, user)
57             correlations[int(user)] = r
58
59     correlations = sorted(correlations.items(), key=lambda x: x
60                          [1])
61     pp(correlations)
62     least_correlated = correlations[:5]
63     most_correlated = correlations[-5:]
64
65     with open("data/correlatedUsers.txt", 'w') as f:
66         print("Most Correlated:", file=f)
67         pp(most_correlated, stream=f)
68         print("Least Correlated:", file=f)
69         pp(least_correlated, stream=f)
70
71 def transformPrefs(prefs):
72     '''
73     Transform the recommendations into a mapping where persons
74     are described
75     with interest scores for a given title e.g. {title: person}
76     instead of
77     {person: title}.
78     '''
79     result = {}
80     for person in prefs:
81         for item in prefs[person]:
82             result.setdefault(item, {})
83             # Flip item and person
84             result[item][person] = prefs[person][item]
85     return result
86
87 def getRecommendations(prefs, person, similarity=sim_pearson):
88     '''
89     Gets recommendations for a person by using a weighted
90     average
91     of every other user's rankings
92     '''

```



```

93     totals = {}
94     simSums = {}
95     for other in prefs:
96         # Don't compare me to myself
97         if other == person:
98             continue
99         sim = similarity(prefs, person, other)
100        # Ignore scores of zero or lower
101        if sim <= 0:
102            continue
103        for item in prefs[other]:
104            # Only score movies I haven't seen yet
105            if item not in prefs[person] or prefs[person][item]
               == 0:
106                # Similarity * Score
107                totals.setdefault(item, 0)
108                # The final score is calculated by multiplying
                   each item by the
109                # similarity and adding these products
                   together
110                totals[item] += prefs[other][item] * sim
111                # Sum of similarities
112                simSums.setdefault(item, 0)
113                simSums[item] += sim
114        # Create the normalized list
115        rankings = [(total / simSums[item], item) for (item, total)
                   in
116                    totals.items()]
117        # Return the sorted list
118        rankings.sort()
119
120        lowestRankings = rankings[:5]
121        topRankings = rankings[-5:]
122
123        return (lowestRankings, topRankings)
124
125
126 def topMatches(
127     prefs,
128     person,
129     n=5,
130     similarity=sim_pearson,
131 ):
132     '''
133     Returns the best matches for person from the prefs
        dictionary.
134     Number of results and similarity function are optional
        params.
135     '''

```

```

136
137     scores = [(similarity(prefs, person, other), other) for
138               other in prefs
139               if other != person]
140     scores.sort()
141     # scores.reverse()
142     lowestScores = scores[:n]
143     highestScores = scores[-n:]
144     return (lowestScores, highestScores)
145
146 def loadMovieLens(path='./'):
147     # Get movie titles
148     movies = {}
149     for line in open(path + 'data/u.item', encoding="ISO
150                     -8859-1"):
151         (id, title) = line.split('|')[0:2]
152         movies[id] = title
153
154     # Load data
155     prefs = {}
156     for line in open(path + 'data/u.data'):
157         (user, movieid, rating, ts) = line.split('\t')
158         prefs.setdefault(user, {})
159         prefs[user][movies[movieid]] = float(rating)
160     return prefs
161
162 def saveScores(filename, lowScores, highScores):
163     with open(filename, 'w') as f:
164         print("Lowest Scores:", file=f)
165         pp(lowScores, stream=f)
166         print("Highest Scores:", file=f)
167         pp(highScores, stream=f)
168
169
170 if __name__ == "__main__":
171     # q2
172     prefs = loadMovieLens()
173     findCorrelations(prefs)
174     # 868 is substituteMe
175     # q3
176     getRecommendations(prefs, '868')
177     # q4
178     prefs = transformPrefs(prefs)
179     (lowestScore, highestScore) = topMatches(prefs, 'Citizen
180         Kane (1941)')
181     saveScores("data/favoriteFilmCorrelation.txt", lowestScore,
182               highestScore)

```

```
181 |     (lowestScore , highestScore) = topMatches(prefs , 'Mars  
    |         Attacks! (1996) ')  
182 |     saveScores("data/worstFilmCorrelation.txt" , lowestScore ,  
    |         highestScore)
```

Listing 2: Python script utilizing Programming Collective Intelligence's code

3

Question

3. Compute ratings for all the films that the substitute you have not seen. Provide a list of the top 5 recommendations for films that the substitute you should see. Provide a list of the bottom 5 recommendations (i.e., films the substitute you is almost certain to hate).

Answer

The answer for this question again relied heavily upon Programming Collective Intelligence’s code since it provided the main method to solve this problem, the *getRecommendations* method, and is shown in Listing 2. This again utilized the Sim Pearson correlation coefficient between users and found all the movies my substitute user, user 868, hasn’t seen. After a final score was calculated these scores were normalized on a 1 to 5 scale and sorted in order. I took the lowest 5 movies and the top 5 movies, again with some being the same weight I simply took the top 5 it provided. These are shown in Table 6 and 7.

I was taken aback when substitute me’s second most recommended movie was “Santa with Muscles (1996).” He also apparently hates any kind of Amityville horror movie. I’ll have to look into Saint of Fort Washington.

Rank	Movie	Rating
1	Saint of Fort Washington, The (1993)	5.0
2	Santa with Muscles (1996)	5.0
3	Someone Else’s America (1995)	5.0
4	The Deadly Cure (1996)	5.0
5	They Made Me a Criminal (1939)	5.0

Table 6: Most recommended movies

Rank	Movie	Rating
1	3 Ninjas: High Noon At Mega Mountain (1998)	1.0
2	Amityville 1992: It's About Time (1992)	1.0
3	Amityville: A New Generation (1993)	1.0
4	Amityville: Dollhouse (1996)	1.0
5	Babyfever (1994)	1.0

Table 7: Least recommended movies

4

Question

4. Choose your (the real you, not the substitute you) favorite and least favorite film from the data. For each film, generate a list of the top 5 most correlated and bottom 5 least correlated films. Based on your knowledge of the resulting films, do you agree with the results? In other words, do you personally like / dislike the resulting films?

Answer

When looked through the movie data I saw one of my favorite movies of all time, Citizen Kane (1941) because of the amazing cinematography for its time, I chose this film as my favorite. For me least favorite film I chose Mars Attacks! (1996). This question used the *getRecommendations* method like in question 3, but this time I had to transform the preferences into a dictionary of movie as key and users as values. The method to do this was *transformPrefs*, which was also provided by the Programming Collective Intelligence book [2]. I then used the *topMatches* method to get the bottom and top most correlated films. The results are shown in Tables 8, 9, 10 and 11 with the data saved to a files named **favoriteFilmCorrelation.txt** and **worstFilmCorrelation.txt** on my Github repository [1].

Its difficult to comment on the results of this problem as I have never seen any of the films in the tables below, aside from Free Willy, but not even that sequel number. I imagine that the films that least correlate with Citizen Kane also align with my dislikes after looking them up on the IMDB website. Aside from that, I really can't comment on the rest of the films.

Rank	Movie	Correlation
1	Newton Boys, The (1998)	1.0
2	Palmetto (1998)	1.0
3	Savage Nights (Nuits fauves, Les) (1992)	1.0
4	Wild America (1997)	1.0
5	Heavy (1995)	1.0

Table 8: Citizen Kane highest correlated movies

Rank	Movie	Correlation
1	Maya Lin: A Strong Clear Vision (1994)	-1.0
2	Free Willy 3: The Rescue (1997)	-1.0
3	Bad Girls (1994)	-1.0
4	Big Bully (1996)	-1.0
5	Colonel Chabert, Le (1994)	-1.0

Table 9: Citizen Kane least correlated movies

Rank	Movie	Correlation
1	Winter Guest, The (1997)	1.0
2	Wonderful, Horrible Life of Leni Riefenstahl, The (1993)	1.0
3	Wooden Man's Bride, The (Wu Kui) (1994)	1.0
4	Kaspar Hauser (1993)	1.0
5	Palmetto (1998)	1.0

Table 10: Mars Attacks! highest correlated movies

Rank	Movie	Correlation
1	Across the Sea of Time (1995)	-1.0
2	8 Heads in a Duffel Bag (1997)	-1.0
3	Swan Princess, The (1994)	-1.0
4	8 Seconds (1994)	-1.0
5	Aparajito (1956)	-1.0

Table 11: Mars Attacks! least correlated movies

5

Question

Extra credit (3 points)

5. Rank the 1,682 movies according to the 1997/1998 MovieLense data. Now rank the same 1,682 movies according to today's (March 2016) IMDB data (break ties based on # of users, for example: 7.2 with 10,000 raters > 7.2 with 9,000 raters).

Draw a graph, where each dot is a film (i.e., 1,682 dots). The x-axis is the MovieLense ranking and the y-axis is today's IMDB ranking.

What is Pearson's r for the two lists (along w/ the p-value)? Assuming the two user bases are interchangeable (which might not be a good assumption), what does this say about the attitudes about the films after nearly 20 years?

Answer

NOT ATTEMPTED

6

Question

Extra credit (3 points)

6. Repeat #6, but IMDB data from approximately July 31, 2005. What is the cumulative error (in days) from the desired target day of July 31, 2005? For example, if 1 memento is from July 1, 2005 and another memento is from July 31, 2006, then the cumulative error for the two mementos is 30 days + 365 days = 385 days.

Note: the URIs in the MovieLens data redirect, be sure to use the final values as URI-Rs for the archives.

Answer

NOT ATTEMPTED

References

- [1] Atkins, Grant. “CS532 Assignment 7 Repository” Github. N.p., 23 March 2017. Web. 23 March 2017.<https://github.com/grantat/cs532-s17/tree/master/assignments/A7>.
- [2] Segaran, Toby. “Programming Collective Intelligence”. O’ Reilly, 2007. Web. 6 April 2017. <http://shop.oreilly.com/product/9780596529321.do>.